



SOFTWARE ENGINEERING PROJECT

**KUru: AI-Powered University Program
Advisor
for KU Program Selection using RAG,
Knowledge Graph, and Hybrid
Recommendation (Proposal)**

BY

**Thanawat Tantijaroensin 6610545294
Phantawat Luengsiriwattana 6610545871**

**DEPARTMENT OF COMPUTER
ENGINEERING
FACULTY OF ENGINEERING
KASETSART UNIVERSITY**

Academic Year 2568

**KUru: AI-Powered University Program
Advisor
for KU Program Selection using RAG,
Knowledge Graph, and Hybrid
Recommendation (Proposal)**

BY

**Thanawat Tantijaroensin 6610545294
Phantawat Luengsiriwattana 6610545871**

**A Project Submitted in Partial Fulfillment of the
Requirement for the Degree of Bachelor of Engineering
(Software and Knowledge Engineering)
Department of Computer Engineering
Faculty of Engineering
KASERTSART UNIVERSITY
Academic Year 2568**

Approved By:

Advisor Date
(Jitti Niramitranon)

Table of Contents

Content	Page
List of Tables	iii
List of Figures	v
Chapter 1 Introduction	1
1.1 Background	1
1.2 Problem Statement	2
1.3 Solution Overview	2
1.3.1 Features	2
1.4 Target Users	5
1.5 Terminology	5
Chapter 2 Literature Review and Related Work	7
2.1 Competitor Analysis	7
2.1.1 KU Official Information Channels	7
2.1.2 General-Purpose LLMs	7
2.1.3 University-Specific AI Advisors at Other Thai In-	
stitutions	8
2.1.4 TCAS Information Platforms	8
2.1.5 Comparison Summary	8
2.2 Literature Review	10
2.2.1 Retrieval-Augmented Generation (RAG)	10
2.2.2 Knowledge Graphs for Educational Recommen-	
dation	10
2.2.3 Adaptive Quiz and Interest Profiling	11
2.2.4 Hybrid Recommendation Architecture	13

Chapter 3	Requirement Analysis	14
3.1	Stakeholder Analysis	14
3.2	User Stories and Use Cases	14
3.2.1	High School Student User Stories	14
3.2.2	Additional User Stories	15
3.2.3	Key Use Cases	15
3.3	Use Case Diagram	26
3.4	Non-Functional Requirements	28
3.5	User Interface Design	29
Chapter 4	Software Architecture Design	36
4.1	System Architecture Overview	36
4.2	Software Design	36
4.2.1	RAG Pipeline Sequence	36
4.2.2	Recommendation Pipeline Sequence	39
4.2.3	Interest Elicitation Flow	41
4.3	AI and Data Design	42
4.3.1	Data Sources and Collection	42
4.3.2	Model Design	45
4.3.3	Evaluation	46
Chapter 5	Software Development	51
5.1	Schedule	51
Appendix A:	RIASEC Elicitation Question Bank and Scoring Logic	55
Appendix B:	About L^AT_EX	59

List of Tables

	Page
1.1 Terminology	6
2.1 Competitor Comparison Summary	9
5.1 Project Schedule	52
A.1 Likert Question Bank (4 statements per RIASEC dimension, 5-point scale, max score 20)	56
A.2 Pairwise Question Bank (6 adjacent dimension pairs)	56
A.3 Scenario Question Bank (3 scenarios $\times$ 6 role options with RIASEC mapping)	57

List of Figures

	Page
3.1 KUru Use Case Diagram	27
3.2 KUru landing page — bilingual entry point with interest-discovery and program-browse paths.	30
3.3 Interest discovery flow: (left) opening screen prompting the student to begin the 24-statement RIASEC assessment; (right) Step 12 profile summary showing the Holland Code, dominant dimensions (Social, Realistic), and per-dimension score bars before proceeding to recommendations.	30
3.4 Recommendation results page — RIASEC profile summary (top), matched O*NET careers with fit scores (middle), and ranked KU program cards (bottom).	31
3.5 Program explorer — grid of ranked program cards with PLO fit scores, faculty tags, and TCAS score ranges. The search bar supports natural language queries in Thai and English.	32
3.6 KUru Advisor chat interface — the student asks “What is CPE?” and receives a grounded answer about Computer Engineering TCAS Round 3 requirements with source citations.	33
3.7 Portfolio Coach: (left) document upload screen accepting portfolio PDFs and transcript images; (right) analysis results showing a 75% readiness score for Computer Engineering, with GPAX and competition award criteria marked as fulfilled.	34
4.1 KUru System Architecture	37
4.2 KUru Domain Model	38

4.3	RAG Pipeline Sequence	48
4.4	Recommendation Pipeline — Five Stages with Two Parallel Signals	49
4.5	Interest Elicitation Flow (UC-01) — 12-Step Adaptive Process	50

Chapter 1

Introduction

1.1 Background

Thailand's university admission system, TCAS (Thai University Central Admission System), requires students to select and apply to specific faculties at specific universities through up to four rounds, each with distinct requirements, quotas, and criteria. Kasetsart University (KU), one of Thailand's leading public universities, offers programs spanning agriculture, engineering, science, humanities, business, and more across its campuses.

Despite the availability of official information about each program, most prospective students make faculty selection decisions with limited structured guidance. Program Learning Outcomes (PLOs), which formally document the competencies a graduate will have developed, are defined for every accredited program in Thailand via the TQF (Thai Qualifications Framework) document system, known as มคอ.2. These documents are the authoritative source of what a program actually teaches and produces, yet they are inaccessible to the students they most directly affect: written in formal Thai, formatted inconsistently across faculties, and distributed as PDFs on faculty websites without any structured interface.

Prospective students who are interested in KU programs — whether finishing high school, taking a gap year, or reconsidering their options — face this problem acutely, as they must make high-stakes decisions with limited structured guidance.

1.2 Problem Statement

High school students selecting KU programs for TCAS lack personalized, grounded guidance that connects their interests to specific program PLOs, career outcomes, and admission requirements. Existing tools are either generic (general-purpose LLMs with no KU-specific data) or static (official faculty web pages with no personalization or synthesis).

The consequence is suboptimal faculty selection, contributing to higher rates of program dissatisfaction, transfer requests, and dropout.

1.3 Solution Overview

KUru is a bilingual (Thai and English) web application that serves as an intelligent advisor for students who are exploring career paths or programs at Kasetsart University. The system is grounded in official KU curriculum documents provided by faculty, making its guidance specific, accurate, and authoritative rather than approximate.

The system conducts adaptive interest discovery, produces ranked program recommendations with PLO visualizations, provides career path detail for matched programs, and presents a structured TCAS admission guide specific to each faculty and round.

The system focuses exclusively on students at the pre-admission stage — before any enrollment decision is made.

The core AI architecture combines a RAG pipeline over มคอ.2 curriculum documents with a Neo4j knowledge graph connecting faculties, PLOs, skills, and careers. All recommendations are grounded in retrieved curriculum content, not generated from model weights alone.

1.3.1 Features

1. **Interest Discovery:** A twelve-step adaptive interest elicitation interface grounded in Holland’s RIASEC vocational model [?]: (1) six sequential Likert rating screens, one per RIASEC dimension (R/I/A/S/E/C), each presenting 4 Thai-language statements rated on a 5-point scale

(1 = ไม่เห็นด้วยอย่างมาก, 5 = เห็นด้วยอย่างมาก) — raw scores per dimension are summed (max 20); (2) one global confidence check screen applying a scalar (1.0 / 0.75 / 0.5) to all six dimension scores; (3) 4–6 adaptive pairwise forced-choice questions targeting only dimension pairs where $|\text{score}_A - \text{score}_B| < 3$ after confidence scaling; (4–6) 3 scenario questions with A–F role options mapped to RIASEC dimensions, adjusting scores proportionally; and (7) a profile summary page with an optional dealbreaker filter zeroing out explicitly rejected dimensions before L2-normalisation. Total elicitation time is 2–4 minutes. As the student accumulates interaction data, a blending weight α shifts the system progressively from RIASEC-based to behavioural-signal-based recommendations.

2. **Program Recommendation Engine:** Ranked list of KU programs matched to the student’s RIASEC interest profile using a hybrid recommendation architecture combining two independent signals: a career-side signal (Pipeline A) matching the student’s RIASEC vector against O*NET occupation data via pgvector, and a curriculum-side signal (Pipeline B) matching against KU program PLO profiles via Neo4j and course content via pgvector semantic search. Final score = $0.35 \times \text{A-score} + 0.65 \times \text{B-score}$. Recommendations are explained in plain language covering the full chain from interests to career alignment to curriculum fit. Covers all KU programs for which มคอ.2 documents are available.
3. **PLO Spider Chart Visualizer:** Interactive radar chart showing the skill profile a student will develop in any given program, overlaid with the student’s interest profile for visual fit assessment.
4. **KUru Advisor:** RAG-powered Q&A interface over มคอ.2 documents supporting natural language queries in Thai and English about any KU program.

5. **TCAS Admission Guide:** Structured per-faculty, per-round admission information including GPAX requirements, exam score criteria (TGAT/TPAT/A-Level), portfolio requirements, and deadlines.
6. **Saved Profile Dashboard:** Optional login to persist interest profile, bookmarked programs, and TCAS deadline tracking across sessions.
7. **PLO Explorer with Semantic Search:** A browsable and searchable directory of all KU programs. Students can search using natural language queries in Thai or English (e.g., “หลักสูตรที่เรียน AI เน้นโปรเจกต์ รับ TGAT ต่ำ”). Queries are parsed into structured constraints covering topic areas, teaching methods, and TCAS requirements, then matched against program content via multi-source parallel query against Neo4j, Supabase TCAS data, and pgvector. Students can pin up to 4 programs to a persistent comparison tray for side-by-side comparison, batch portfolio checking, and chatbot queries with multi-program context.
8. **Portfolio Readiness Checker:** Students upload a portfolio PDF; Gemini extracts a structured portfolio profile (activities, certificates, academic records, personal statement). The system compares this against pre-extracted faculty criteria schemas for selected programs and performs a four-part gap analysis: hard threshold eligibility check, required item coverage, preferred item strength assessment, and qualitative criteria assessment. Output is a structured report with eligibility status, portfolio strength profile, prioritised gap list, and deadline-aware action recommendations. MVP scope: top 10 programs. Excludes programs requiring creative work evaluation.
9. **Curriculum Timeline Visualiser:** For any KU program, displays a visual map of what the student will study across 4 years. Course composition, teaching method breakdown

(lecture/lab/project ratios), and time commitment per year are extracted from มคอ.2 course structure data during ingestion. Gemini synthesises this structured data into natural-language year narratives on demand. This is structured extraction followed by grounded generation — not open-ended RAG.

10. **Program Comparison:** Students can pin up to 4 programs from the explorer or recommendation screen and compare them side-by-side, showing: overlaid PLO radar charts, career path overlap, curriculum character profile across 5 dimensions (theory vs. project, individual vs. team, lab vs. lecture, early specialisation vs. broad foundation, industry connection vs. academic focus), TCAS accessibility per round, and personalised fit scores relative to the student's RIASEC profile.

1.4 Target Users

KUru serves prospective students at the pre-admission stage:

- **Prospective KU applicants (pre-admission):** Anyone considering KU programs before making an enrollment decision. This includes high school students in Mathayom 4–6 preparing for TCAS, gap-year students reconsidering their options, and anyone exploring what KU programs offer and what careers they lead to. No prior KU enrollment is assumed. Guidance spans from initial interest discovery through to understanding specific TCAS admission requirements.

Users may access the system without an account. Login unlocks profile persistence, saved programs, and deadline notifications.

1.5 Terminology

Table 1.1: Terminology

Term	Definition
PLO	Program Learning Outcome — a formally defined competency that graduates of a specific program are expected to have developed, as documented in มคอ.2.
CLO	Course Learning Outcome — a learning outcome defined at the individual course level, documented in มคอ.3. CLOs contribute to PLOs.
มคอ.2	The Thai Qualifications Framework program specification document for each accredited university program. Contains PLOs, curriculum structure, and admission requirements.
มคอ.3	The course-level specification document detailing CLOs, teaching methods, and assessment for each course.
TCAS	Thai University Central Admission System — the national university admission platform through which students apply to Thai universities in up to four rounds per year.
TGAT/TPAT/A-Level	Standardized national examinations used as admission criteria in TCAS Rounds 2–4.
RAG	Retrieval-Augmented Generation — an AI architecture that retrieves relevant passages from a document corpus and provides them as context to a language model before generating a response, grounding output in retrieved evidence.
Knowledge Graph	A graph database structure in which nodes represent entities (faculties, PLOs, skills, careers) and edges represent relationships between them (a faculty produces certain PLOs, a PLO develops certain skills, a career requires certain skills).
O*NET	The Occupational Information Network — a US Department of Labor database cataloguing approximately 900 occupations with structured skill requirements. Used as the career-to-skill data layer.
RIASEC	Holland’s six vocational personality types: Realistic, Investigative, Artistic, Social, Enterprising, Conventional. Used as the theoretical framework for interest elicitation and as the native classification system of the O*NET occupational database.
Holland Code	A two-letter code formed from a person’s two most dominant RIASEC types (e.g., IR for Investigative-Realistic), used as a human-readable summary of an individual’s vocational interest profile.

Chapter 2

Literature Review and Related Work

2.1 Competitor Analysis

Five categories of existing solutions are relevant to this project:

2.1.1 KU Official Information Channels

KU's official website and faculty pages provide static information about programs, curricula, and admission requirements. PLO information is available only within มคอ.2 PDF documents, which are distributed without structured interfaces. TCAS information is scattered across the main KU admissions page, individual faculty pages, and the national mytcas.com platform. No personalization, synthesis, or guidance is offered. The fundamental limitation is that information exists but is not actionable for students making decisions.

2.1.2 General-Purpose LLMs

General-purpose language models such as ChatGPT and Gemini can respond to natural language queries about KU programs, but do so from training data rather than authoritative sources. This results in approximations and hallucinations when asked about specific PLOs, current TCAS requirements, or quota numbers. These models have no persistent memory of the student across sessions and provide no structured outputs such as visualizations or deadline trackers.

2.1.3 University-Specific AI Advisors at Other Thai Institutions

Knowva (KMITL) and ChulaGENIE (Chulalongkorn University) represent the closest competitors to KUru. Both are AI-assisted advising systems built for specific Thai universities. However, neither is available at KU, neither explicitly links PLOs to career skill requirements, and neither provides TCAS-specific admission guidance. ChulaGENIE focuses on course recommendation for enrolled students rather than pre-admission guidance.

2.1.4 TCAS Information Platforms

Dedicated TCAS information platforms, most prominently mytcas.com (the official national TCAS portal operated by the Ministry of Higher Education), aggregate admission data across all Thai universities into a single structured interface. For each program, mytcas.com presents per-round quotas, eligibility criteria (curriculum type, GPAX thresholds), score calculation methods (portfolio weighting, TGAT/TPAT/A-Level proportions), and links to faculty admission pages. The data is updated each TCAS cycle and is the authoritative source consulted by both students and universities.

Despite its comprehensiveness, mytcas.com is a passive information browser, not a guidance system. It presents raw admission data without connecting it to a student's interests, skills, or career goals. Students must independently navigate from a program they have already identified to its admission requirements; the platform provides no discovery mechanism, no program matching, and no PLO or career-path context. KUru is designed to complement mytcas.com by providing the personalized discovery and program-matching layer that sits upstream of the admission details mytcas.com supplies.

2.1.5 Comparison Summary

Table 2.1 summarises how KUru addresses the gaps identified across the five categories of existing solutions.

Table 2.1: Competitor Comparison Summary

Feature	KU Web-site	mytcas .com	ChatGPT/ Gemini	Knowva/ ChulaGE-NIE	KUru (ours)
KU-specific curriculum data	Partial (PDF)	✗	✗	✗	Full (มคอ.2 RAG)
PLO visualization	✗	✗	✗	✗	Spider chart
Career-to-PLO mapping	✗	✗	~	~	Full graph
TCAS admission guide	Scattered	Full — all uni-versities	~	✗	Structured, KU only
Personalization	✗	✗	Session-only	Basic	Progressive profile
Thai language support	✓	✓	~	~	✓
KU student skill tracking	✗	✗	✗	✗	Yes (Phase 2)

2.2 Literature Review

2.2.1 Retrieval-Augmented Generation (RAG)

Lewis et al. [?] introduced RAG as an architecture combining dense retrieval with sequence-to-sequence generation for open-domain question answering. The key insight is that grounding generation in retrieved evidence substantially reduces hallucination and allows the model to answer questions requiring specific factual knowledge not reliably encoded in model weights. This is directly applicable to KUru’s requirement to answer questions about specific PLO content, admission criteria, and course structures from มคอ.2 documents.

For evaluation, the RAGAS framework [?] provides automated metrics for RAG pipeline quality including faithfulness, answer relevancy, and context recall, enabling objective assessment of whether the system’s answers are grounded in retrieved curriculum content.

2.2.2 Knowledge Graphs for Educational Recommendation

Knowledge graph-based recommendation systems have been demonstrated in educational contexts to enable multi-hop reasoning that purely embedding-based approaches cannot support. Embedding-based retrieval can identify semantically similar items but cannot follow explicit typed relationships such as ‘this PLO develops this skill’ or ‘this career requires this skill level.’ The Neo4j graph in KUru encodes these typed relationships explicitly, enabling recommendation paths that are both explainable and traceable to specific curriculum documents [?, ?].

Using knowledge graphs as context sources for LLM-generated explanations has been specifically investigated by Abu-Rasheed et al. [?], who propose a conversational explainability approach where knowledge graph data is injected into LLM prompts to reduce hallucination risk while maintaining natural dialogue. This is the precise architectural pattern KUru uses for its Program Recommendation Engine: Neo4j provides the structured PLO-to-skill data, and Gemini generates the natural-language explanation from that data rather than from general knowledge.

2.2.3 Adaptive Quiz and Interest Profiling

KUru’s interest discovery interface is grounded in Holland’s RIASEC model (Realistic, Investigative, Artistic, Social, Enterprising, Conventional), the dominant framework in vocational psychology for connecting individual interests to occupational environments [?]. This theoretical foundation is directly compatible with KUru’s O*NET career data layer [?], which uses RIASEC codes to classify all 900+ occupations. Rather than presenting abstract RIASEC labels, KUru presents six sequential Likert screens using relatable Thai-language statements grounded in Holland’s six dimensions, reducing cognitive load for a high school student audience while producing a more reliable initial vector than a tile-selection approach. The elicitation design is defined in Section 4.3.2.

The interface design is shaped by three evidence-based principles from the preference elicitation (PE) and recommender systems literature.

First, the system uses explicit attribute-based elicitation rather than item-based elicitation. Research on large-scale onboarding systems demonstrates that item-based approaches — asking users to evaluate individual items directly — are unsuitable when the item corpus is large, because no single item efficiently maps a new user’s preference space [?]. KUru instead presents topic clusters (attributes) that map directly to RIASEC dimensions and downstream skill vectors, enabling broad preference coverage in minimal interactions.

Second, the questionnaire is deliberately minimal. Sepiarskaia et al. [?] demonstrated at ACM RecSys that an optimized Static Preference Questionnaire (SPQ) can reduce necessary questionnaire length by up to a factor of three compared to standard elicitation methods, while maintaining or improving recommendation quality. While the SPQ principle advocates for minimalism, KUru’s five-step design remains substantially shorter than validated instruments while incorporating adaptive question selection — targeting only ambiguous RIASEC dimensions rather than administering all questions uniformly — which preserves the efficiency principle while improving signal quality. KUru’s twelve-step adaptive design — (1) six Likert screens (4 statements per dimension, 5-point scale); (2) one global confidence check; (3) 4–6 adaptive pairwise ques-

tions targeting only ambiguous dimension pairs ($\text{delta} < 3$ points); (4) three scenario questions confirming the dominant signal; (5) an optional dealbreaker filter on the profile summary page — is substantially shorter than a full validated interest inventory while producing a more reliable initial vector than a single tile-selection approach. The pairwise comparison step is adaptive: students with a clearly differentiated profile encounter fewer questions, keeping total completion time within 2–4 minutes.

Third, the system addresses the new user cold-start problem [?] through this explicit elicitation phase, then transitions to progressive implicit profiling. Behavioral signals including program page visits, chatbot queries, and saved items continuously refine the user’s RIASEC vector without additional questioning [?, ?]. The structured grid interface is supported by de Vries et al. [?], who found in a controlled user study that high-guidance elicitation — presenting structured options rather than open-ended prompts — produces significantly higher recommendation match scores than low-guidance alternatives.

It should be noted that interest-based elicitation has inherent limitations as a predictor of long-term academic fit. Meta-analytic research on full validated interest inventories reports approximately 50% accuracy in predicting eventual career choice even with substantially longer instruments [?]. KUru’s abbreviated elicitation is accordingly framed as a starting point for exploration rather than a definitive match, with recommendation confidence increasing as the implicit profiling layer accumulates behavioral signals over subsequent sessions.

This work applies established cold-start PE methods and the RIASEC vocational framework to academic pathway advising for Thai university applicants — a novel application context in which the user population (Thai high school students navigating TCAS), the item space (university programs defined by PLOs), and the decision stakes (a multi-year academic commitment) differ substantially from the content recommendation domains in which these methods were originally developed.

2.2.4 Hybrid Recommendation Architecture

KUru's recommendation pipeline is a hybrid system combining two independent signals: a knowledge-based signal derived from matching the student's RIASEC vector against O*NET occupation profiles (Pipeline A), and a content-based signal derived from matching against KU program PLO profiles and course content via Neo4j and pgvector semantic search (Pipeline B). Hybrid recommender systems that combine independent evidence sources consistently outperform single-signal approaches by reducing the error compounding that occurs when one signal feeds sequentially into the next [?]. In the previous single-chain design, an approximation introduced at the O*NET-to-SkillCluster crosswalk step would propagate through to the final ranking; the parallel architecture eliminates this dependency. The final score weights curriculum-side evidence more heavily (65%) than career-side evidence (35%), reflecting the system's primary purpose: helping students understand what they will actually study at KU, not merely what careers they might eventually pursue.

Chapter 3

Requirement Analysis

3.1 Stakeholder Analysis

- **High school students (prospective KU applicants):** Primary users. Need personalized program recommendations, PLO-based skill previews, career path information, and structured TCAS guidance. Low technical proficiency expected; interface must be accessible and engaging.
- **KU Faculty / Academic Advisors:** Data providers. Provide มคอ.2 curriculum documents and TCAS admission data. Indirect beneficiaries through reduced advising load and better-prepared incoming students.
- **KU Admissions Office:** Indirect stakeholders. Benefit from better-informed applicants and reduced mismatched enrollments.
- **Development Team:** Two SKE students responsible for design, development, and evaluation within one semester.

3.2 User Stories and Use Cases

3.2.1 High School Student User Stories

- As a high school student, I want to discover which KU programs match my interests so that I can make an informed TCAS application.

- As a high school student, I want to see what skills I will graduate with from a specific program so that I can assess whether it aligns with my career goals.
- As a high school student, I want to understand the TCAS admission requirements for a program I am interested in so that I know what I need to prepare.
- As a high school student, I want to ask questions about a program's curriculum in Thai so that I can get specific answers without reading long PDF documents.
- As a high school student, I want to save programs I am considering so that I can return to compare them later.

3.2.2 Additional User Stories

- As a returning visitor, I want my saved RIASEC interest profile and bookmarked programs to be available when I come back so that I do not have to start over.
- As a parent, I want to browse KU programs and understand what careers they lead to so that I can support my child's decision.

3.2.3 Key Use Cases

UC-01: Discover Interests

Actor: Guest / High School Student

Precondition: None — accessible without login.

Main Flow:

1. Student opens KUru and navigates to ค้นหาความสนใจ. System presents six sequential screens, one per RIASEC dimension (R, I, A, S, E, C), each showing 4 Thai-language statements rated on a 5-point Likert scale (1 = ไม่เห็นด้วยอย่างมาก, 5 = เห็นด้วยอย่างมาก). A progress bar shows current step out of 12 total.

2. System computes a raw score per dimension (sum of 4 responses, max 20). Before proceeding to the pairwise step, system presents one confidence check screen: “คุณรู้สึกมั่นใจแค่ไหนกับคำตอบที่เพิ่งเลือก?” with three options (มั่นใจมาก = scalar 1.0, ค่อนข้างมั่นใจ = 0.75, ไม่แน่ใจเลย = 0.5). The scalar is applied globally to all six Likert-derived dimension scores.
3. System identifies ambiguous dimension pairs — any two dimensions where $|\text{score}_A - \text{score}_B| < 3$ after confidence scaling. Only pairs meeting this threshold are shown as pairwise forced-choice questions (maximum 4–6 pairs). Students with a clearly differentiated profile receive fewer questions. Each pair presents two options with “ชอบเท่ากัน” as a middle choice.
4. System presents 3 scenario questions (สถานการณ์ที่ 1–3), each offering A–F role options mapped to RIASEC dimensions. Student selects one preferred role per scenario. Scenario responses adjust dimension scores proportionally.
5. System displays the profile summary page showing the top 2 dominant dimensions, a bar chart of all 6 dimension scores (out of 20), and a Holland Code label. Below the summary, system presents an optional dealbreaker filter: “มีสาขาไหนที่คุณไม่อยากจะเรียนเลย?” showing 6 dimension chips. Student may tap any to exclude; tapped dimensions are zeroed in the final vector regardless of score. A “ไม่มีซ้ำมาได้เลย” option is prominently shown.
6. System applies L2-normalisation to the final 6-dimensional vector and stores it as the student’s RIASEC profile. Student proceeds to recommendations via “ดูผลลัพธ์”.

Alternative Flow — Targeted re-elicitation: If invoked from UC-02 with specific RIASEC dimensions flagged for clarification, the system skips Steps 1–2 and begins directly at Step 3, presenting only pairwise questions

targeting the flagged dimensions. Steps 4–6 run normally. This path is triggered when the student selects “No” on the profile correction prompt in UC-02.

Postcondition: Student has a RIASEC-grounded interest profile used as input to the career matching and program recommendation pipeline. Total elicitation time is 2–4 minutes across 12 steps.

UC-02: View Program Recommendations

Actor: Guest / High School Student

Precondition: Interest profile (RIASEC vector) has been built via UC-01.

Main Flow:

1. System runs Pipeline A: queries precomputed O*NET occupation RIASEC interest profiles in Supabase pgvector for occupations with the highest cosine similarity to the student’s RIASEC vector. Top 10–15 occupations form the internal career space. An A-score is computed per KU program based on career-program pathway alignment.
2. System runs Pipeline B concurrently: (B1) queries Neo4j for KU programs whose PLO SkillCluster profiles best match the student’s RIASEC-derived SkillCluster weight vector via cosine similarity; (B2) runs semantic search over course description embeddings in Supabase pgvector using the student’s interest query. $B\text{-score} = 0.4 \times B1 + 0.6 \times B2$.
3. System synthesises a final score: $0.35 \times A\text{-score} + 0.65 \times B\text{-score}$. Programs are ranked by final score. Top- N programs selected for display.
4. System displays ranked program cards, each showing fit score, matched career paths from Pipeline A, curriculum alignment themes from Pipeline B, and a plain-language explanation generated by Gemini grounded in both signals.
5. Student can tap any card to navigate to the full program detail page at `/programs/[program-id]` — a full-page view, not a modal.

6. System presents a profile correction prompt — “Do these recommendations feel right?” — on the first visit to the recommendation screen. If the student selects “No”, the system returns them to a targeted re-elicitation of the RIASEC dimensions most influential in the current ranking, preventing a poor initial vector from entering the interaction log.

Postcondition: Student sees a ranked list of KU programs with recommendations traceable from their interests through real-world careers to specific program outcomes via two independent evidence paths.

UC-03: Explore PLO Spider Chart

Actor: Guest / High School Student

Precondition: A program card is selected.

Main Flow:

1. Student selects a KU program from the recommendation list.
2. System retrieves PLO data from Supabase and skill mappings from Neo4j.
3. System renders an interactive radar (spider) chart showing the skill profile the program develops.
4. Student’s interest profile is overlaid on the chart for visual fit comparison.

Postcondition: Student has a visual understanding of the skills a program builds and how well its PLO skill profile aligns with their RIASEC interest vector.

UC-04: Query KUru Advisor

Actor: Guest / High School Student

Precondition: ๓๓๐.2 documents have been ingested and indexed in Supabase pgvector.

Main Flow:

1. Student enters a question in Thai or English in the chat interface.

2. System embeds the query using Gemini text-embedding-001.
3. System retrieves the top semantically relevant มคอ.2 document chunks from pgvector.
4. System passes the chunks and question to Gemini 2.5 Flash with a strict citation instruction.
5. System displays the generated answer with inline มคอ.2 source citation badges.
6. The chatbot also handles TCAS questions for any program where admission data has been ingested — for example, “Computer Engineering Round 1 ต้องการ GPAX เท่าไหร่?” or “Architecture portfolio ต้องมีอะไรบ้าง?” Answers are grounded in the ingested TCAS data and include source citations. If TCAS data for a requested program has not been ingested, the system returns the standard fallback message.

Alternative Flow: If no relevant chunk is found, the system responds: “This information was not found in the curriculum document.”

Postcondition: Student has received a cited answer grounded in the official curriculum document.

UC-05: View TCAS Admission Guide

Actor: Guest / Registered User

Precondition: TCAS admission data has been collected and structured per faculty and round. Students may also query TCAS information conversationally via KUru Advisor (UC-04), which draws from the same ingested data source.

Main Flow:

1. Student navigates to the TCAS Guide for a faculty they are interested in.
2. System displays a structured breakdown of applicable rounds, score requirements (GPAX, TGAT/TPAT/A-Level), portfolio criteria, and deadlines.

3. Student can filter by round and view side-by-side comparisons of faculties.

Postcondition: Student understands the specific admission requirements they need to prepare for.

UC-06: Explore Career Paths

Actor: Guest / High School Student

Precondition: RIASEC interest profile has been built via UC-01.

Main Flow:

1. Student taps “ดูอาชีพที่เหมาะสมกับคุณ” (View matching careers) from the recommendation screen.
2. System retrieves the top 7 O*NET occupations (selected from the 10–15 matched internally during the recommendation pipeline) to display in the career explorer.
3. System displays a career list; each card shows the occupation title, a Thai-localised description, the top 3 required skills, and the KU programs that develop those skills.
4. Student taps a career card to expand detail.
5. System shows the full O*NET skill breakdown (35 standardized skills) and highlights which matched KU programs develop each skill, using Neo4j data.
6. Student can navigate directly from a career card to the relevant program card.

Postcondition: Student understands which careers align with their RIASEC interest profile and which KU programs prepare them for those careers.

UC-07: Refine Profile Through Implicit Signals

Actor: Registered User

Precondition: Student is logged in and has an existing RIASEC profile from UC-01.

Main Flow:

1. Student browses program cards, saves programs, checks TCAS guides, and asks chatbot questions across one or more sessions.
2. System records behavioral signals alongside the PLO skill profiles of the associated programs, weighted by interaction type:
 - Program save — weight 1.0
 - TCAS guide viewed — weight 0.8
 - Chatbot query about a program — weight 0.6
 - Program card expanded — weight 0.4
 - Program card viewed — weight 0.1
3. System adjusts a blending weight α based on accumulated interaction data. A new user starts at $\alpha = 1.0$ (pure RIASEC-based recommendations). As interaction weight accumulates across 3–5 meaningful signals, α decays toward 0.2, shifting the ranking progressively toward behavioural signal. The transition is smooth and invisible to the student. At each session, the recommendation ranking is computed as:

$$\text{score} = \alpha \times \text{RIASEC fit score} + (1 - \alpha) \times \text{behavioural fit score}$$
4. On the next session, the recommendation ranking reflects the updated blended profile.
5. Student can view a summary of how their recommendation ranking has shifted under “โปรไฟล์ของคุณ” on the saved profile dashboard, including a before/after comparison of the highest-ranked programs.

Alternative Flow: Student disagrees with the ranking shift and taps “รีเซ็ตโปรไฟล์” to clear accumulated interaction data and return to pure RIASEC recommendations ($\alpha = 1.0$).

Note on guest users: Interaction signals are only logged persistently for registered users. Guest interactions are tracked within the current session

only and do not update α across visits. Guests always begin a new session at $\alpha = 1.0$ (pure RIASEC recommendations).

Postcondition: Recommendation ranking becomes progressively more personalised to the student’s actual behaviour without requiring additional questionnaire steps. The underlying RIASEC vector from elicitation remains fixed; the behavioural fit score is derived from the individual student’s own interaction history (not from other users) and operates as a separate signal blended at ranking time via α . True collaborative filtering — using patterns from students with similar profiles — is deferred to Phase 2 (see Section ??).

UC-08: Explain Why a Program Was Recommended

Actor: Guest / High School Student

Precondition: Student is viewing a program recommendation card.

Main Flow:

1. Student taps “ทำไมถึงแนะนำคณะนี้?” (Why was this recommended?) on any program card.
2. System retrieves the full reasoning chain from both pipelines: Pipeline A — student RIASEC vector $\rightarrow$ matched O*NET occupations $\rightarrow$ career alignment score; Pipeline B — student RIASEC vector $\rightarrow$ SkillCluster weights $\rightarrow$ Neo4j PLO match + course content semantic match. Both signals contributed to this program’s ranking.
3. System generates a plain-language explanation via Gemini referencing both signals: “Based on your interest in [topic], careers like [Y] and [Z] align with your profile. This program was also strongly matched based on its curriculum content — specifically courses in [X] and [Y] which directly address your interest areas.”
4. Student can tap any career name in the explanation to navigate to its full detail in UC-06.

Postcondition: Student understands the specific chain of reasoning behind a recommendation and can verify each step themselves.

UC-09: Search Programs Semantically

Actor: Guest / High School Student

Precondition: None — accessible without login.

Main Flow:

1. Student enters a natural language query in Thai or English in the semantic search bar on the PLO Explorer (e.g., “หลักสูตรที่เรียน AI เน้นโปรเจกต์ รับ TGAT ต่ำ”).
2. System uses Gemini to parse the query into a structured constraint object: topic areas, teaching methods, TCAS requirements, campus preferences.
3. System executes a multi-source parallel query: Neo4j for topic/skill constraints, Supabase for TCAS constraints, pgvector for semantic content matching against course descriptions.
4. System returns ranked program cards matching the constraints, with an explanation of which constraints each program satisfies.
5. Student can refine the query or apply faceted filters (faculty, program level, TCAS round).

Postcondition: Student finds programs matching specific criteria without browsing the full catalogue.

UC-10: Pin Programs and Compare

Actor: Guest / High School Student

Precondition: At least 2 programs pinned via the pin button on any program card.

Main Flow:

1. Student pins programs from the explorer or recommendation screen. Pinned programs appear in a persistent tray at the top of the explorer (maximum 4).
2. Student taps “เปรียบเทียบ” from the pin tray.

3. System generates a side-by-side comparison: overlaid PLO radar charts, career path overlap, curriculum character profiles across 5 dimensions, TCAS accessibility per round, and personalised fit scores relative to the student's RIASEC profile.
4. Student can tap “เช็ค Portfolio ทั้งหมด” to run a portfolio gap analysis against all pinned programs simultaneously, showing which program the student is most ready for right now.
5. Student can tap “แชทกับ KUru” to open the chatbot with all pinned programs as active context.

Postcondition: Student has a clear comparative picture of their shortlisted programs and knows which they are most ready to apply for.

UC-11: View Curriculum Timeline

Actor: Guest / High School Student

Precondition: Student is on a program detail page.

Main Flow:

1. Student selects the “ชีวิตนักศึกษา” tab on a program detail page.
2. System retrieves course structure data extracted from มคอ.2 during ingestion: courses per year, credit hour distribution, teaching method breakdown (lecture/lab/project ratios), assessment method types, and sample learning activity descriptions.
3. System displays a year selector (ปี 1–4). Default shows Year 1.
4. For the selected year, system displays a Gemini-generated narrative of what that year involves, course composition breakdown, workload indicator, and one concrete example activity description extracted from มคอ.2.
5. Student taps different years to compare progression across the program.

Postcondition: Student has a concrete sense of what studying this program year by year actually involves before applying.

UC-12: Portfolio Readiness Check

Actor: Guest / High School Student

Precondition: Student has a portfolio PDF. Target program has portfolio criteria data ingested.

Main Flow:

1. Student navigates to the Portfolio Coach from the nav bar, a program detail page, or the pin tray.
2. Student selects target program(s) (MVP: from top 10 supported programs) and uploads their portfolio PDF.
3. System uses Gemini to extract a structured portfolio profile: activities classified by type and level, GPAX, certificates, personal statement presence.
4. System retrieves the pre-extracted faculty criteria schema for the selected program(s).
5. System performs a four-part gap analysis: (1) hard threshold eligibility check; (2) required item coverage; (3) preferred item strength assessment with level-aware scoring; (4) qualitative criteria assessment using multi-step Gemini reasoning.
6. System displays a structured report: eligibility status, portfolio strength profile, prioritised gap list, and deadline-aware action recommendations.
7. If multiple programs are selected, system shows which program the student is most ready for right now.

Postcondition: Student knows exactly what is missing from their portfolio and what to prioritise before the application deadline. The system provides preparation guidance, not admission probability prediction.

3.3 Use Case Diagram




Figure 3.1: Kuru Use Case Diagram

3.4 Non-Functional Requirements

- **Performance:** The RAG chatbot must return a grounded response within 5 seconds for 95% of queries under normal load. The recommendation pipeline must return ranked results within 8 seconds of interest profile submission. The PLO spider chart must render within 2 seconds of program selection.
- **Availability:** The web application must maintain at least 95% uptime during the TCAS application period (January–April). Planned maintenance must be performed outside this window.
- **Scalability:** The system must support at least 100 concurrent users during peak TCAS periods without degradation of the above response time targets. Supabase connection pooling and Vercel’s serverless deployment model are relied upon to meet this target without custom infrastructure.
- **Security:** All user authentication is delegated to Supabase Auth (OAuth 2.0 / magic link). Portfolio PDFs uploaded for the Portfolio Readiness Checker are processed ephemerally via the Gemini API and are not retained on the server after processing. API keys for Gemini and Neo4j are stored as environment variables and never exposed to the client.
- **Data Privacy (PDPA):** Portfolio PDFs and interest profiles contain personal data. In accordance with Thailand’s Personal Data Protection Act B.E. 2562 (PDPA), users are notified of data usage at the point of upload. Guest interaction data is session-scoped and not persisted. Registered users may delete their stored profile and interaction history at any time via the dashboard.
- **Accessibility:** The interface must meet WCAG 2.1 Level AA contrast requirements. All interactive elements must have accessible labels for screen reader compatibility. Thai

and English language switching must be available on every screen.

- **Maintainability:** TCAS admission data and มคอ.2 curriculum documents must be re-ingestible independently of the application codebase. The data ingestion pipeline must produce a full re-index without manual code changes when new PDF documents are supplied.

3.5 User Interface Design

The interface follows a mobile-first web design using Next.js, Tailwind CSS, and Shadcn/UI. The key screens are described below.

Landing Page. The entry point presents a bilingual Thai/English value proposition with two primary call-to-action paths: ค้นหาความสนใจ (interest discovery) for undecided students, and ดูโปรแกรม 70+ คณะ มก. (browse programs) for students who already have a direction. Supporting copy highlights the three-minute completion time, มคอ.2 data coverage, and free access.

Interest Discovery. The twelve-step adaptive elicitation flow collects the student’s RIASEC profile. Steps 1–6 present one Likert screen per RIASEC dimension (4 Thai-language statements, 5-point scale, colour-coded per dimension); Step 7 applies a global confidence scalar; Step 8 shows adaptive pairwise forced-choice questions for ambiguous dimension pairs (maximum 4–6, threshold $\delta < 3$); Steps 9–11 present three scenario questions with A–F role options; Step 12 displays the profile summary with Holland Code, per-dimension score bars, and optional deal-breaker filter chips. A progress bar is shown throughout.

Recommendation Results (ผลลัพธ์). After elicitation, the system presents a personalised results page structured in three sections: (1) the



Figure 3.2: KURu landing page — bilingual entry point with interest-discovery and program-browse paths.




Figure 3.3: Interest discovery flow: (left) opening screen prompting the student to begin the 24-statement RIASEC assessment; (right) Step 12 profile summary showing the Holland Code, dominant dimensions (Social, Realistic), and per-dimension score bars before proceeding to recommendations.

student’s RIASEC profile card displaying their Holland Code and top dimensions; (2) matched O*NET occupations via Pipeline A, each showing a fit percentage, required skills, and linked KU programs; (3) ranked KU program cards with fit scores derived from Pipeline B. On first visit, a profile correction prompt allows the student to trigger targeted re-elicitation if the ranking does not feel right.




Figure 3.4: Recommendation results page — RIASEC profile summary (top), matched O*NET careers with fit scores (middle), and ranked KU program cards (bottom).

Program Explorer (สำรวจหลักสูตร). A browsable and semantically searchable directory of KU programs. Each program card displays the PLO fit score, faculty tag, estimated TCAS score range, and entry-round badges. Students can filter by faculty or keyword and pin up to four programs to a persistent comparison tray at the top of the page.

KUru Advisor (แชทกับ KUru). A RAG-powered conversational interface for querying curriculum and TCAS information. The chatbot answers natural language questions in Thai and English, grounding responses in ingested มคอ.2 documents and TCAS admission data. Each




Figure 3.5: Program explorer — grid of ranked program cards with PLO fit scores, faculty tags, and TCAS score ranges. The search bar supports natural language queries in Thai and English.

response includes source citations (document name and page). When programs are pinned, the chatbot supports multi-program context queries.




Figure 3.6: KUru Advisor chat interface — the student asks “What is CPE?” and receives a grounded answer about Computer Engineering TCAS Round 3 requirements with source citations.

Portfolio Coach. A two-step portfolio readiness checker accessible from the navigation bar, program detail page, or pin tray. In the first step the student uploads portfolio documents (PDF or image). In the second step the system displays a readiness score, a gap analysis comparing the student’s credentials against the target program’s admission criteria, and a checklist of fulfilled and missing items.

Program Detail Page. Full-page view at `/programs/[program-id]` — not a modal — to support URL sharing, extended reading, and mobile scrolling. Contains: hero section with program identity, career section




Figure 3.7: Portfolio Coach: (left) document upload screen accepting portfolio PDFs and transcript images; (right) analysis results showing a 75% readiness score for Computer Engineering, with GPAX and competition award criteria marked as fulfilled.

(matched occupations from Pipeline A), curriculum timeline tab (year-by-year visualiser via UC-11), PLO radar chart overlaid with the student's interest profile, TCAS section with personal eligibility check, and portfolio checker entry point (UC-12).

Program Comparison View. Side-by-side comparison triggered from the pin tray (UC-10). Shows overlaid PLO radar charts, career path overlap, curriculum character profile across five dimensions, TCAS accessibility per round, and RIASEC fit scores for up to four pinned programs simultaneously.

Career Explorer. List of O*NET occupations matched to the student's RIASEC profile via Pipeline A. Each card shows the Thai-localised occupation title, required skill clusters, and the KU programs that develop those skills.

TCAS Guide. Faculty-specific breakdown of applicable admission rounds, score requirements (GPAX, TGAT, TPAT, A-Level), portfolio criteria, and application deadlines, drawn from the ingested TCAS dataset.

Saved Profile Dashboard. Bookmarked programs, TCAS deadline tracker, interest profile summary, and a before/after comparison of recommendation ranking shifts over time. Requires student login.

Chapter 4

Software Architecture Design

4.1 System Architecture Overview

KUru is a three-tier web application with a distinct AI layer. The frontend (Next.js) communicates with a Python FastAPI backend, which orchestrates two AI subsystems: a RAG pipeline for curriculum Q&A and a graph-based recommendation engine. Both subsystems read from a shared data layer consisting of Supabase (PostgreSQL + pgvector for embeddings) and Neo4j (knowledge graph on Railway).

Figure 4.2 shows the domain model: the eleven core entities and their relationships. A **Student** holds one **RIASECProfile** (the normalised 6-dimensional interest vector produced by UC-01) and generates **InteractionLog** entries as they browse. Each **Program** belongs to a **Faculty**, which defines its **PLOs** and publishes **PortfolioCriteria**. PLOs are linked to **SkillCluster** nodes in Neo4j (the *develops* relationship), and **Career** nodes are linked to the same clusters (the *requires* relationship) — this shared vocabulary enables the multi-hop recommendation path from student interests to program outcomes. **TCASRecord** and **Course** are extracted structured layers stored in Supabase alongside the vector embeddings.

4.2 Software Design

4.2.1 RAG Pipeline Sequence

When a user submits a query to the curriculum chatbot: (1) the query is embedded using Gemini text-embedding-001; (2) the embedding is used to perform a cosine similarity search against the pgvector store of




Figure 4.1: Kuru System Architecture




Figure 4.2: Kuru Domain Model

๓๓๐.2 document chunks; (3) the top- k retrieved chunks are assembled into a prompt context along with the user query and the conversation history; (4) Gemini 2.5 Flash generates a grounded response in the requested language; (5) source citations are appended to the response.

4.2.2 Recommendation Pipeline Sequence

When a student completes the interest discovery flow, the recommendation pipeline proceeds in five stages using two independent parallel signals:

1. **RIASEC vector construction.** The twelve-step adaptive elicitation (UC-01) produces a weighted 6-dimensional RIASEC vector. Steps 1–6 (Likert screens) produce raw scores per dimension (max 20 each). Step 7 applies a global confidence scalar (1.0, 0.75, or 0.5). Step 8 (adaptive pairwise) adjusts scores for dimension pairs with $\Delta < 3$. Steps 9–11 (scenario questions) adjust scores proportionally per response. Step 12 applies the optional dealbreaker filter, zeroing out explicitly rejected dimensions. The resulting vector is L2-normalised to a unit vector before use in Pipeline A and Pipeline B.
2. **Parallel signal computation.** Two independent pipelines run concurrently:

Pipeline A (Career-side): The student RIASEC vector is compared against precomputed O*NET occupation RIASEC interest profiles in Supabase pgvector using cosine similarity. O*NET interest profiles are empirically derived from surveys of employed workers, making this step grounded in real-world data. The top 10–15 occupations form the internal career space; the top 7 are surfaced to the student in the Career Explorer (UC-06). An A-score per KU program is computed based on career-program pathway alignment. O*NET’s 35-skill taxonomy is used for Career

Explorer display in UC-06 only — it is not used as a bridge to Neo4j in the ranking pipeline.

Pipeline B (Curriculum-side): Two sub-signals computed independently and combined:

- **B1 — Neo4j PLO match:** The student’s RIASEC vector is converted to a SkillCluster weight vector and matched against KU program PLO SkillCluster profiles via Neo4j graph query. B1-score is the cosine similarity between the student weight vector and each program’s PLO skill profile.
- **B2 — Semantic course content match:** The student’s interest query is matched against course description embeddings in Supabase pgvector via semantic search over ingested มคอ.2 course content. B2-score is the weighted mean similarity of the top retrieved course chunks per program.

Combined B-score = $0.4 \times B1 + 0.6 \times B2$.

The two pipelines are independent — an error in Pipeline A does not propagate into Pipeline B.

3. **Score synthesis.** Programs are ranked by a final score combining both pipeline outputs:

$$\text{score} = 0.35 \times \text{A-score} + 0.65 \times \text{B-score}$$

The top- N programs are selected for display.

4. **Explanation generation.** The top- N ranked programs, together with their Pipeline A details (matched O*NET occupations) and Pipeline B details (curriculum themes and matched PLOs), are passed to Gemini 2.5 Flash as structured context. Gemini generates a plain-language explanation for each program referencing both career alignment and curriculum fit. Explanations are generated in Thai or English depending on the student’s language setting.

5. **Behavioural re-ranking (registered users only).** For registered users with accumulated interaction data, the pipeline scores from Stage 3 are blended with a behavioural fit score derived from the individual student's own interaction history (not from other users — true collaborative filtering across students is deferred to Phase 2). The blend ratio is controlled by α : new users start at $\alpha = 1.0$ (pure pipeline score); α decays toward 0.2 as interaction weight accumulates according to the threshold schedule documented in UC-07 (3–5 meaningful signals $\rightarrow \alpha \approx 0.5$; regular user $\rightarrow \alpha \approx 0.2$). The final ranking score is:

$$\text{final score} = \alpha \times \text{pipeline score} + (1 - \alpha) \times \text{behavioural fit}$$

Guest interaction signals are session-scoped and do not update α ; guests always begin a new session at $\alpha = 1.0$. Fallback hierarchy when behavioural data is sparse: (1) sufficient interaction history $\rightarrow$ normal blended signal; (2) limited interaction history $\rightarrow \alpha$ held higher; (3) no interaction history $\rightarrow$ pure pipeline score ($\alpha = 1.0$). Interaction weights are documented in UC-07.

4.2.3 Interest Elicitation Flow

The interest elicitation flow (UC-01) produces the RIASEC vector that seeds both the recommendation pipeline and the career explorer. The flow is adaptive: the pairwise comparison step (Step 8) is shown only when at least one dimension pair has a score delta below the 3-point ambiguity threshold; students with a clearly differentiated Likert profile skip directly to the scenario questions. Figure 4.5 shows the full conditional logic from entry to RIASEC vector output.

4.3 AI and Data Design

4.3.1 Data Sources and Collection

- **KU curriculum documents (มคอ.2):** Provided directly by KU faculty advisors for all KU programs. Thai-language PDFs of varying length (typically 20–80 pages per program). These are the primary knowledge base for both the RAG pipeline and the PLO extraction pipeline. Documents are version-tagged by academic year and re-ingested when updated.
- **O*NET occupational database [?]:** Publicly available from the US Department of Labor. Contains two data layers used by KUru: (a) *Interests* — the RIASEC profile for each occupation, derived from surveys of employed workers, used for career matching from the student’s RIASEC vector; and (b) *Skills* — 35 standardized skills per occupation with importance and level scores, displayed in the Career Explorer (UC-06) to show students what skills their matched careers require. Not used as a bridge to Neo4j in the ranking pipeline — Pipeline B1 queries Neo4j directly from the RIASEC-derived SkillCluster weight vector. Downloaded as structured data (CSV/JSON). Occupation titles are mapped to Thai equivalents for student-facing display; skill and interest data are used directly as the underlying structure is language-independent.
- **TCAS admission data:** Ingested from three sources in priority order. (1) *KU faculty-provided project PDFs* received via Google Drive — each PDF (e.g., โครงการข้างเฟือก.pdf) may contain admission records for multiple programs across multiple faculties. A batch downloader retrieves all PDFs recursively. Each PDF is classified as tcas / mko2 / portfolio / unknown using rule-based keyword detection on the first page. TCAS PDFs are processed by a structured extraction pipeline using Gemini to parse per-program

TCASRecord objects covering project name, round, faculty, program name, quota, GPAX minimum, exam score criteria, portfolio requirements, and application deadlines. (2) *mytcas.com* — scraped as a gap-filling layer for programs not covered by faculty PDFs, using the `mytcas_code` field in the canonical program registry. (3) *Faculty portfolio criteria PDFs* — processed separately using Gemini structured extraction to produce per-faculty, per-round criteria schemas for the Portfolio Readiness Checker. All three sources are resolved against a canonical program registry using fuzzy name matching with PyThaiNLP normalization and sequence similarity scoring (threshold 0.85). Records below this threshold are flagged for manual review. A Supabase materialized view (`programs_unified`) joins มคอ.2 structured data, TCAS records, and portfolio criteria into a unified program record keyed on `program_id`, refreshed after each ingestion run.

- **Interaction log:** Generated by the system as registered users engage with KUru. Stored in Supabase PostgreSQL. Records program saves, card expansions, TCAS guide views, chatbot queries, and card views, each with a timestamp, interaction weight, and associated program identifier. Used to compute the behavioural fit score in Stage 5 of the recommendation pipeline. A background recomputation job updates collaborative scores periodically as new interaction data accumulates.
- **Faculty portfolio criteria documents:** Published by KU faculties for TCAS Round 1 and Round 2 applications. Ingested via the same PDF extraction pipeline as มคอ.2 using Gemini in structured extraction mode to produce a per-faculty, per-round criteria schema. Stored in Supabase PostgreSQL alongside TCAS admission data. Updated at the start of each TCAS cycle. Used by the Portfolio Readiness Checker (Feature 8).

- **Course structure data (extracted layer):** Extracted from มคอ.2 during ingestion as a structured layer stored separately from the embedding chunks. Fields extracted per program per year: course list, credit hour distribution, teaching method breakdown (lecture/lab/project ratios), assessment method types, and sample learning activity descriptions. Stored in Supabase PostgreSQL. Used by the Curriculum Timeline Visualiser (Feature 9).

มคอ.2 document ingestion uses a six-stage pipeline:

1. **PDF classification.** Pages are classified by text density and Thai character density (minimum 20% Thai Unicode characters for Thai-language documents, minimum 100 characters per page) into born-digital, scanned, pure-image, or mixed types.
2. **Tiered text extraction.** Born-digital pages are processed by PyMuPDF directly. Pages failing Thai character density or length checks are re-processed using Typhoon OCR (opentyphoon.ai), a vision-language model purpose-built for Thai and English document extraction that outperforms general-purpose VLMs on Thai government document formats including complex table layouts and mixed Thai/English content. Gemini Vision serves as a secondary fallback if the Typhoon OCR API is unavailable. Pages where both fallbacks fail are flagged for manual review. Each extracted page is tagged with its extraction method (pymupdf / typhoon_ocr / gemini_vision / failed) and confidence level for quality monitoring.
3. **Section-aware chunking.** PLO sections, course descriptions, admission requirements, and curriculum structure sections are chunked separately (target 500 tokens, 50-token overlap) to preserve semantic coherence. Each chunk is tagged with section type (plo / course / admission / curriculum_structure / general).

4. **Structured extraction.** PLOs are extracted as typed JSON objects (plo_id, description_th, description_en, domain, skill_keywords) for Neo4j population. Curriculum structure is extracted as a per-year, per-semester JSON object for the Curriculum Timeline Visualiser.
5. **Embedding generation.** Gemini text-embedding-001 generates embeddings in batches with rate limiting. Chunks from failed OCR pages are excluded.
6. **Neo4j population and quality reporting.** Program → PLO → SkillCluster nodes and edges are created from the structured PLO extraction output. An ingestion report is generated per program flagging documents where OCR ratio exceeds 50% or PLO extraction returns zero results for manual review.

4.3.2 Model Design

The primary language model is Gemini 2.5 Flash, selected for its strong Thai language performance, native PDF document understanding, 1M token context window, and cost efficiency relative to GPT-4o. Direct Gemini API calls are used without LangChain or similar orchestration frameworks, to reduce latency and maintain explicit control over prompt structure.

Embeddings use Gemini text-embedding-001, a multilingual model that produces consistent semantic representations for both Thai and English text, enabling cross-language retrieval where a Thai query can retrieve an English document chunk and vice versa.

The Neo4j knowledge graph uses a schema with four node types: Faculty, PLO, SkillCluster, and Career. Edges represent: Faculty -[HAS_PLO]→ PLO, PLO -[DEVELOPS]→ SkillCluster, Career -[REQUIRES]→ SkillCluster. This schema supports the multi-hop queries required for recommendation: given a student's SkillCluster profile, find Faculties whose PLOs develop those clusters.

The interest discovery interface uses a 12-step elicitation flow grounded in Holland’s RIASEC model [?]. Steps 1–6 present one RIASEC dimension per screen with 4 Likert statements each (5-point scale, max score 20 per dimension). Step 7 applies a global confidence scalar ($1.0 / 0.75 / 0.5$) to all dimension scores based on a single self-reported confidence question. Step 8 presents adaptive pairwise forced-choice questions targeting only dimension pairs where the score delta is below a 3-point threshold, with a maximum of 4–6 pairs shown. Steps 9–11 present three scenario questions with A–F role options mapped to RIASEC dimensions, adjusting scores proportionally. Step 12 displays the profile summary with an optional dealbreaker filter that zeros out any explicitly rejected dimensions before L2-normalisation. The full question bank is documented in Appendix A.

The recommendation ranking blends the RIASEC-derived fit score with a behavioural fit score using a blending weight α that decays automatically as the student accumulates interaction data: new users start at $\alpha = 1.0$ (pure RIASEC); after 3–5 meaningful interactions $\alpha \approx 0.5$ (balanced blend); regular users converge toward $\alpha \approx 0.2$ (behaviour-dominated). Interaction signals are weighted by type: program save (1.0), TCAS guide viewed (0.8), chatbot query about a program (0.6), card expanded (0.4), card viewed (0.1).

The baseline for comparison is BM25 keyword retrieval over the same ๓๓๐.2 corpus without embedding or LLM generation. The full RIASEC-to-topic mapping is documented in Appendix A.

4.3.3 Evaluation

- **RAG pipeline quality:** RAGAS framework metrics [?] — Faithfulness (proportion of claims in the answer supported by retrieved context), Answer Relevancy (semantic similarity between answer and question), Context Recall (proportion of relevant information retrieved). Target: Faithfulness > 0.80 , Answer Relevancy > 0.75 .
- **Recommendation quality:** Mean Reciprocal Rank (MRR) and NDCG@5 evaluated on a manually curated test set of

30–50 student interest profiles with ground-truth faculty matches validated by faculty advisors or senior students. Target: $\text{MRR} > 0.60$.

- **User experience:** Task completion rate and System Usability Scale (SUS) score from user testing with 10–15 high school students or KU first-year students. Target: SUS score > 70 (acceptable).




Figure 4.3: RAG Pipeline Sequence




Figure 4.4: Recommendation Pipeline — Five Stages with Two Parallel Signals




Figure 4.5: Interest Elicitation Flow (UC-01) — 12-Step Adaptive Process

Chapter 5

Software Development

5.1 Schedule

Development is organized into four phases across April 2026 and January March 2027. P1 Foundation runs April 1–12, P2 Core AI runs April 6–30, P3 Features runs January 2027, and P4 Polish & Evaluation runs February–March 2027. By the end of April 2026, core AI features (RAG chatbot and recommendation engine) will be completed as a proof of concept.

Table 5.1: Project Schedule

Period	Phase	Deliverables
Apr 1–5, 2026	P1 · Foundation — Set up Repo	Repository setup, project structure, development environment
Apr 1–5, 2026	P1 · Foundation — Data Collection	Receive มคอ.2 and TCAS admission data from KU faculty, download O*NET dataset
Apr 6–12, 2026	P1 · Foundation — Data Ingestion	PDF extraction, chunking, embedding into Supabase pgvector; Neo4j schema + pilot data
Apr 6–30, 2026	P2 · Core AI — RAG Chatbot PoC	RAG pipeline (retrieval + Gemini generation), TCAS Q&A support
Apr 6–30, 2026	P2 · Core AI — RIASEC Recommendation PoC	Recommendation engine (Pipeline A + B, interest elicitation)
Apr 6–30, 2026	P2 · Core AI — Basic Integration & First Demo	End-to-end integration of RAG + Recommendation, first working demo
Jan 1–31, 2027	P3 · Features — Program Explorer	PLO Explorer with semantic search and pin tray
Jan 1–31, 2027	P3 · Features — Program Detail Pages	Full program detail pages, curriculum timeline visualiser, PLO spider chart
Jan 1–31, 2027	P3 · Features — Portfolio Coach	Portfolio readiness checker (top 10 programs)
Jan 31, 2027	P3 · Features — Complete Recommendation	Complete recommendation feature integration
Feb 1–28, 2027	P4 · Polish & Eval — UI/UX Polish	UI/UX polish, Thai/English language switching, mobile-first refinement
Feb 1–28, 2027	P4 · Polish & Eval — Cross-feature Integration	Chatbot + Explorer + Portfolio connection
Feb 1–28, 2027	P4 · Polish & Eval — User Testing	RAGAS evaluation, MRR/NDCG test set, SUS user testing with 10–15 students
Mar 1–31, 2027	P4 · Polish & Eval — Performance Testing	Performance benchmarking against all success metric targets
Mar 1–31, 2027	P4 · Polish & Eval — Bug Fixing & Optimization	Bug fixes, optimization, stability improvements
Mar 1–31, 2027	P4 · Polish & Eval — Final Documentation	Final report, demo preparation, project documentation

Reference

Appendix A

Appendix A: RIASEC Elicitation Question Bank and Scoring Logic

This appendix documents the full question bank used by the 12-step interest discovery interface (UC-01) and the scoring rules applied at each stage. After all steps are complete, the 6-dimensional RIASEC vector is L2-normalised before use in the recommendation pipeline.

Adaptive threshold note: Pairwise questions (Table A.2) are shown only for dimension pairs where $|\text{score}_A - \text{score}_B| < 3$ after confidence scaling. A maximum of 4–6 pairs are presented per session.

Table A.1: Likert Question Bank (4 statements per RIASEC dimension, 5-point scale, max score 20)

Dimension	Q	Statement (Thai)
Realistic (R)	1	ฉันชอบงานที่ได้ใช้มือและร่างกาย เช่น ประกอบ ซ่อม หรือสร้างสิ่งของ
	2	ฉันสนุกกับการทำงานกับเครื่องมือ เครื่องจักร หรืออุปกรณ์ต่าง ๆ
	3	ฉันชอบแก้ปัญหาที่จับต้องได้มากกว่าการคิดเชิงนามธรรม
	4	ฉันรู้สึกสบายใจเมื่อได้ทำงานกลางแจ้งหรือในพื้นที่จริง
Investigative (I)	1	ฉันชอบวิเคราะห์ข้อมูลและหาคำตอบจากหลักฐาน
	2	ฉันสนุกกับการตั้งสมมติฐานและทดสอบแนวคิดใหม่ ๆ
	3	ฉันชอบอ่านหรือค้นคว้าเรื่องที่ฉันสนใจอย่างลึกซึ้ง
	4	ฉันรู้สึกสนุกเมื่อได้แก้ปัญหาที่ซับซ้อนหรือต้องใช้ตรรกะ
Artistic (A)	1	ฉันชอบแสดงความคิดสร้างสรรค์ผ่านงานศิลปะ ดนตรี หรือการเขียน
	2	ฉันมักมองเห็นโอกาสในการออกแบบหรือปรับแต่งสิ่งต่าง ๆ ให้สวยงามขึ้น
	3	ฉันรู้สึกสบายใจในสภาพแวดล้อมที่เปิดโอกาสให้แสดงตัวตน
	4	ฉันชอบงานที่ไม่มีกฎตายตัวและมีอิสระในการสร้างสรรค์
Social (S)	1	ฉันรู้สึกมีพลังเมื่อได้ช่วยเหลือหรืออธิบายให้คนอื่นเข้าใจ
	2	ฉันชอบทำงานเป็นทีมและสร้างความสัมพันธ์กับผู้อื่น
	3	ฉันสนใจปัญหาสังคมและอยากมีส่วนร่วมในการแก้ไข
	4	ฉันชอบให้คำปรึกษาหรือสนับสนุนเพื่อน ๆ เมื่อพวกเขาต้องการ
Enterprising (E)	1	ฉันชอบโน้มน้าวคนอื่นและนำเสนอแนวคิดของตัวเอง
	2	ฉันสนุกกับการวางแผนโครงการและนำทีมให้บรรลุเป้าหมาย
	3	ฉันรู้สึกตื่นเต้นกับความท้าทายและโอกาสทางธุรกิจ
	4	ฉันชอบตัดสินใจและรับผิดชอบผลลัพธ์ที่เกิดขึ้น
Conventional (C)	1	ฉันชอบงานที่มีขั้นตอนชัดเจนและทำตามระบบได้
	2	ฉันรู้สึกสบายใจเมื่องานมีความแม่นยำและตรวจสอบได้
	3	ฉันชอบจัดระเบียบข้อมูล แฟ้ม หรือตารางให้เป็นระบบ
	4	ฉันทำงานได้ดีในสภาพแวดล้อมที่มีกฎและความคาดหวังที่ชัดเจน

Table A.2: Pairwise Question Bank (6 adjacent dimension pairs)

Pair	Option A	Option B
R vs. I	ออกแบบและสร้างเครื่องมือทางกายภาพ	วิเคราะห์ข้อมูลและพัฒนาทฤษฎีใหม่
I vs. A	ค้นคว้าและทดสอบสมมติฐานทางวิทยาศาสตร์	สร้างสรรค์ผลงานศิลปะหรือสื่อสารสร้างสรรค์
A vs. S	ออกแบบประสบการณ์หรือผลงานที่สวยงาม	สอน ให้คำปรึกษา หรือดูแลผู้อื่น
S vs. E	ทำงานสนับสนุนและช่วยเหลือชุมชน	นำทีมและขับเคลื่อนองค์กรให้เติบโต
E vs. C	บริหารโครงการและสร้างธุรกิจใหม่	จัดการระบบและตรวจสอบความถูกต้อง
R vs. C	งานภาคสนามที่ต้องใช้ทักษะเชิงปฏิบัติ	งานสำนักงานที่เน้นความเป็นระบบและความแม่นยำ

Table A.3: Scenario Question Bank (3 scenarios $\times$ 6 role options with RIASEC mapping)

Scenario	Setting	Role option	RIASEC
1	งานแยกการรอน 48 ชั่วโมง	A — ออกแบบและสร้าง prototype ที่ใช้งานได้จริง B — วิเคราะห์ข้อมูลและหาโซลูชันที่ดีที่สุด C — ออกแบบ UI/UX ให้น่าใช้งานและสวยงาม D — ประสานงานทีมและดูแลสมาชิกให้ทำงานได้ดี E — นำเสนอผลงานต่อกรรมการและโน้มน้าวให้ชนะ F — จัดการตารางงาน งบประมาณ และเอกสารทีม	R I A S E C
2	โปรเจกต์ชุมชน	A — ก่อสร้างหรือซ่อมแซมโครงสร้างพื้นฐาน B — ศึกษาปัญหาและออกแบบแนวทางแก้ไขเชิงระบบ C — ผลิตสื่อสร้างสรรค์เพื่อสร้างการรับรู้ D — พูดคุยกับชาวบ้านและรวบรวมความต้องการ E — ระดมทุนและหาพาร์ทเนอร์สนับสนุนโครงการ F — ติดตามความคืบหน้าและรายงานผลให้ถูกต้อง	R I A S E C
3	สตาร์ทอัพด้านสุขภาพ	A — พัฒนาอุปกรณ์หรือระบบฮาร์ดแวร์ที่ใช้กันได้ B — วิจัยและทดสอบประสิทธิภาพของผลิตภัณฑ์ C — สร้างแบรนด์และออกแบบประสบการณ์ผู้ใช้ D — ฝึกอบรมและสนับสนุนผู้ใช้งานในชุมชน E — วางกลยุทธ์ธุรกิจและขยายตลาด F — ดูแลการปฏิบัติตามกฎระเบียบและมาตรฐาน	R I A S E C

Appendix B

Appendix B: About L^AT_EX

LaTeX (stylized as L^AT_EX) is a software system for typesetting documents. LaTeX markup describes the content and layout of the document, as opposed to the formatted text found in WYSIWYG word processors like Google Docs, LibreOffice Writer, and Microsoft Word. The writer uses markup tagging conventions to define the general structure of a document, to stylize text throughout a document (such as bold and italics), and to add citations and cross-references.

LaTeX is widely used in academia for the communication and publication of scientific documents and technical note-taking in many fields, owing partially to its support for complex mathematical notation. It also has a prominent role in the preparation and publication of books and articles that contain complex multilingual materials, such as Arabic and Greek.

Overleaf has also provided a 30-minute guide on how you can get started on using L^AT_EX. [?]